

Retail

September 4, 2022

1 Importing dataset and libraries

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[30]: df=pd.read_csv('test_data_hidden.csv')
# df= pd.read_excel('train.xlsx')
df.head()
```

```
[30]:
```

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	StateHoliday	\
0	1	5	2015-07-31	5263	555	1	1	0	
1	2	5	2015-07-31	6064	625	1	1	0	
2	3	5	2015-07-31	8314	821	1	1	0	
3	4	5	2015-07-31	13995	1498	1	1	0	
4	5	5	2015-07-31	4822	559	1	1	0	

	SchoolHoliday
0	1
1	1
2	1
3	1
4	1

2 EDA and Visualization

```
[31]: # shape of dataset
print('Shape of dataset')
print('*'*30)
print(f'rows : {df.shape[0]} ')
print(f'columns : {df.shape[1]}')
```

Shape of dataset

rows : 34565

columns : 9

```
[32]: # info
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34565 entries, 0 to 34564
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Store            34565 non-null  int64
1   DayOfWeek        34565 non-null  int64
2   Date             34565 non-null  object
3   Sales            34565 non-null  int64
4   Customers        34565 non-null  int64
5   Open             34565 non-null  int64
6   Promo            34565 non-null  int64
7   StateHoliday     34565 non-null  int64
8   SchoolHoliday    34565 non-null  int64
dtypes: int64(8), object(1)
memory usage: 2.4+ MB
```

```
[33]: # nan values in dataset
df.isna().sum()
```

```
[33]: Store            0
      DayOfWeek       0
      Date            0
      Sales           0
      Customers       0
      Open            0
      Promo           0
      StateHoliday    0
      SchoolHoliday   0
      dtype: int64
```

```
[34]: # no missing values
```

```
[35]: #unique values
df['StateHoliday'].unique()
```

```
[35]: array([0])
```

```
[36]: sh={'0':0, 'a':1, 'b':2, 'c':3,0:0}
```

```
[37]: df['StateHoliday']=df['StateHoliday'].map(sh)
```

```
[38]: df['SchoolHoliday'].unique()
```

```
[38]: array([1, 0])
```

```
[39]: # statistical info
df.describe().T
```

```
[39]:
```

	count	mean	std	min	25%	50%	75%	\
Store	34565.0	558.000000	321.877302	1.0	279.0	558.0	837.0	
DayOfWeek	34565.0	4.000000	1.917688	1.0	2.0	4.0	6.0	
Sales	34565.0	6142.705511	3606.356960	0.0	4325.0	6085.0	8063.0	
Customers	34565.0	643.827224	435.207851	0.0	445.0	610.0	812.0	
Open	34565.0	0.873369	0.332564	0.0	1.0	1.0	1.0	
Promo	34565.0	0.419355	0.493461	0.0	0.0	0.0	1.0	
StateHoliday	34565.0	0.000000	0.000000	0.0	0.0	0.0	0.0	
SchoolHoliday	34565.0	0.369651	0.482717	0.0	0.0	0.0	1.0	

	max
Store	1115.0
DayOfWeek	7.0
Sales	32547.0
Customers	4783.0
Open	1.0
Promo	1.0
StateHoliday	0.0
SchoolHoliday	1.0

```
[40]: df.head()
```

```
[40]:
```


	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	StateHoliday	\
0	1	5	2015-07-31	5263	555	1	1	0	
1	2	5	2015-07-31	6064	625	1	1	0	
2	3	5	2015-07-31	8314	821	1	1	0	
3	4	5	2015-07-31	13995	1498	1	1	0	
4	5	5	2015-07-31	4822	559	1	1	0	

	SchoolHoliday
0	1
1	1
2	1
3	1
4	1

```
[41]: # distribution plot
sns.distplot(df['Sales'])
```

```
/usr/local/lib/python3.9/site-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
warnings.warn(msg, FutureWarning)
```


```
[41]: <AxesSubplot:xlabel='Sales', ylabel='Density'>
```



```
[42]: sns.distplot(df['Customers'])
```

```
/usr/local/lib/python3.9/site-packages/seaborn/distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed in a  
future version. Please adapt your code to use either `displot` (a figure-level  
function with similar flexibility) or `histplot` (an axes-level function for  
histograms).  
warnings.warn(msg, FutureWarning)
```


```
[42]: <AxesSubplot:xlabel='Customers', ylabel='Density'>
```



```
[43]: # boxplots
sns.boxplot(df['Sales'])
```

```
/usr/local/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
  warnings.warn(
```

```
[43]: <AxesSubplot:xlabel='Sales'>
```



```
[44]: #feature engineering
df['day']=pd.to_datetime(df['Date'], format='%Y-%m-%d').dt.day
df['month']=pd.to_datetime(df['Date'], format='%Y-%m-%d').dt.month
df['year']=pd.to_datetime(df['Date'], format='%Y-%m-%d').dt.year
```

```
[45]: df.head()
```

```
[45]:
```

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	StateHoliday	\
0	1	5	2015-07-31	5263	555	1	1	0	
1	2	5	2015-07-31	6064	625	1	1	0	
2	3	5	2015-07-31	8314	821	1	1	0	
3	4	5	2015-07-31	13995	1498	1	1	0	
4	5	5	2015-07-31	4822	559	1	1	0	

	SchoolHoliday	day	month	year
0	1	31	7	2015
1	1	31	7	2015
2	1	31	7	2015
3	1	31	7	2015
4	1	31	7	2015

```
[46]: df.tail()
```

```
[46]:
```

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	\
34560	1111	3	2015-07-01	3701	351	1	1	
34561	1112	3	2015-07-01	10620	716	1	1	
34562	1113	3	2015-07-01	8222	770	1	1	
34563	1114	3	2015-07-01	27071	3788	1	1	
34564	1115	3	2015-07-01	7701	447	1	1	

	StateHoliday	SchoolHoliday	day	month	year
34560	0	1	1	7	2015
34561	0	1	1	7	2015
34562	0	0	1	7	2015
34563	0	0	1	7	2015
34564	0	0	1	7	2015

```
[47]: df.drop('Date', axis=1, inplace=True)
```

```
[48]: df.head()
```

```
[48]:
```

	Store	DayOfWeek	Sales	Customers	Open	Promo	StateHoliday	\
0	1	5	5263	555	1	1	0	
1	2	5	6064	625	1	1	0	
2	3	5	8314	821	1	1	0	
3	4	5	13995	1498	1	1	0	
4	5	5	4822	559	1	1	0	

	SchoolHoliday	day	month	year
0	1	31	7	2015
1	1	31	7	2015
2	1	31	7	2015
3	1	31	7	2015
4	1	31	7	2015

```
[49]: df['StateHoliday'].unique()
```

```
[49]: array([0])
```

```
[50]: y=df['Sales']
X=df.drop('Sales', axis=1)
```

3 Model Building

```
[51]: from sklearn.model_selection import cross_val_score, train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,
↳random_state=42)
```

```
[52]: # linear regression
from sklearn.linear_model import LinearRegression
```

```

lr=LinearRegression()
lr.fit(X_train,y_train)
pred_lr=lr.predict(X_test)
score_lr=cross_val_score(lr,X,y,cv=5)
print(score_lr)

```

[0.8272264 0.74292577 0.79467188 0.8325215 0.82853658]

```
[53]: score_lr.mean()
```

[53]: 0.8051764243690333

```

[54]: from sklearn.metrics import mean_absolute_error, mean_squared_error
mae_lr=mean_absolute_error(y_test,pred_lr)
print(mae_lr)
mse_lr=mean_squared_error(y_test,pred_lr)
print(mae_lr)

```

1020.9858020361147
1020.9858020361147

```

[55]: from sklearn.tree import DecisionTreeRegressor
dt=DecisionTreeRegressor()
dt.fit(X_train,y_train)
pred_dt=dt.predict(X_test)
score_dt=cross_val_score(dt,X,y,cv=5)
print(score_dt)

```

[0.79298482 0.72923396 0.78229266 0.81046471 0.8135094]

```
[56]: score_dt.mean()
```

[56]: 0.785697108855009

```

[57]: mae_dt=mean_absolute_error(y_test,pred_dt)
print(mae_dt)
mse_dt=mean_squared_error(y_test,pred_dt)
print(mae_dt)

```

1031.1549255026762
1031.1549255026762

```

[71]: from sklearn.ensemble import RandomForestRegressor
rf=RandomForestRegressor()
rf.fit(X_train,y_train)
pred_rf=rf.predict(X_test)
score_rf=cross_val_score(rf,X,y,cv=4)
print(score_rf)

```

[0.86766729 0.90337174 0.89618014 0.89128587]


```
[72]: score_rf.mean()
```

```
[72]: 0.8896262596783888
```

```
[73]: mae_rf=mean_absolute_error(y_test,pred_rf)
      print(mae_rf)
      mse_rf=mean_squared_error(y_test,pred_rf)
      print(mae_rf)
```

```
816.019401128309
```

```
816.019401128309
```

```
[74]: # best model is Random forest regressor with accuracy with greater than 89%
```

```
[ ]:
```